

MPI: Project - Game of Life

Yanwen Xu

June 2020

Abstract

In this project, I have written the 1D column decomposition(Pass level) and the 2D tile decomposition(Giga version) and Parallel I/O (Peta version).

1 PCAM

I went through the following PCAM process when designing my code:

- To start, I **Partitioned** the Game-of-Life(GoF) in domain. I decomposed the computation and data into small tasks in two ways: *1D Column Decomposition* and *2D Decomposition*. For 1D decomposition, I split the 2D GoF grids to columns. For 2D decomposition, I split the GoF grids into rectangle tiles.
- For 1D decomposition, I created asynchronous **Communication** channels for each small task to send and receive data(Ghost cells) from its left and right neighbor. For 2D decomposition, each small task will need to send/receive data from all of its 8 neighbor tiles.
- For **Agglomeration**, I evaluated my task-channel communication structures of both 1D and 2D. To reduce necessary communication, I only send the ghost cells (like the left/right most column in 1D) rather than the entire board to other processor. For more comprehensive evaluation of 1D vs. 2D, see Section 4. I used *MPI.Scatter* and *MPI.Gather* to distribute smaller task to process from the main thread.
- For **Mapping** tasks to processors: due to the limited time in development, I chose it to be completely static, and **assume** the dimension of the GoF grids is divisible by the number of processors. So that I can easily assign equal amount of smaller grids to each processor.

2 Documentation

For the full detail of the specifications, please refer to the **README.md** that came along with the project.

3 Proof of Work

To prove that my GoF works correctly, I setup a 20x20 board, with the initial pattern look like in Figure 1. I ran 80 steps, and the resulting pattern has returned to its starting locations as shown in the Figure. Also, in Figure 2, I have shown the result of running 1 step and 4 step. Notice at 4 step that this returns back to the same pattern but with everything shifted diagonally by 1 square.

For the parallel IO, I used *MPI.FILE.WRITE* to allow each processor to write to the same file in parallel. To do so I have to reserve bytes in the file, and allocate each rank with its own region in the file so that there is no when conflict. I have logged the state of the GoF simulation at each of the following step: {0, 19, 39, 59, 79}. Therefore, in the file you can see what each processors contains.

